

Product Requirements Document: ServiceNow Platform Assistant

Author: [Your Name] **Date:** March 6, 2026 **Status:** Draft — Pending Approval **Version:** 1.0

Executive Summary

The **ServiceNow Platform Assistant** is an AI-powered interactive troubleshooting and configuration tool for ServiceNow Solutions Consultants. It provides a conversational interface — similar to developer tools like Claude Code or GitHub Copilot — that enables SCs to explore, diagnose, and configure ServiceNow instances through natural language.

The critical differentiator: All AI/LLM processing runs entirely within the ServiceNow instance using the customer's corporate-approved or BYOK (Bring Your Own Key) LLM providers via the GenAI Controller. **No customer data ever leaves the ServiceNow platform.** This makes it safe for POC environments with real customer data — a scenario where external AI tools like Claude Code or ChatGPT cannot be used due to data privacy policies.

Problem Statement

Current Pain Points

1. **Privacy barrier to AI-assisted work:** SCs cannot use modern AI coding assistants (Claude Code, Copilot, ChatGPT) when working on POC instances containing customer data. Our data privacy policies prohibit sending customer data to external AI services.
2. **Steep learning curve for instance exploration:** New and experienced SCs spend significant time navigating the ServiceNow platform — querying tables, debugging configurations, understanding ACLs, tracing business rules — through manual UI clicks and script debugging.
3. **Repetitive troubleshooting patterns:** Common issues (broken ACLs, misconfigured flows, missing records, performance bottlenecks) follow predictable patterns, yet every SC must rediscover solutions independently.
4. **Setup complexity:** Configuring ServiceNow features (AI Agent Use Cases, Skills Kit skills, integrations) requires creating numerous records across multiple tables in a specific order. Mistakes are silent and hard to diagnose.

The Opportunity

By routing AI capabilities through ServiceNow's own GenAI Controller, we can build a tool that gives SCs the power of AI-assisted troubleshooting and configuration **without any**

data privacy concerns. The tool uses whatever LLM the customer has already approved
— Now LLM, Azure OpenAI (BYOK), AWS Bedrock, or any other configured provider.

Proposed Solution


What It Is

A conversational AI assistant that connects to a ServiceNow instance and helps SCs:

- **Explore:** Query tables, inspect records, understand schemas, check configurations
- **Diagnose:** Find errors in logs, check ACLs, trace business rule execution, identify misconfigurations
- **Configure:** Create and update records, deploy configurations, set up features

How It Works

The assistant runs as a **server-side agentic loop** on the ServiceNow instance itself. The SC interacts through a lightweight command-line interface (CLI) that sends questions and receives answers. All data processing, LLM reasoning, and tool execution happens within the ServiceNow platform boundary.



Key Capabilities

Capability	Examples
Natural language queries	“Show me all P1 incidents from the last week”
Configuration diagnosis	“Why aren’t my assignment rules working?”
Schema exploration	“What fields does the change_request table have?”
Log analysis	“Show me recent errors related to the CMDB”
Setup assistance	“Create a business rule that sets priority based on impact and urgency”
Platform knowledge	“How does the ACL evaluation order work in ServiceNow?”

Target Users

User	Role	Primary Use Case
Solutions Consultants	Primary	POC setup, troubleshooting, demos
Technical Consultants	Secondary	Implementation debugging, configuration
Platform Architects	Secondary	Architecture exploration, best practice validation

Data Privacy & Security

This is the core value proposition. The architecture ensures:

Concern	How It's Addressed
Customer data exposure	All data stays within the ServiceNow instance. The CLI only sends questions (no customer data) and receives sanitized answers.
LLM provider compliance	Uses whatever LLM the customer has already approved via GenAI Controller — Now LLM, BYOK Azure OpenAI, AWS Bedrock, etc.
Access control	All queries use GlideRecordSecure, which respects the authenticated user's ACLs and role-based access.
Audit trail	All configuration changes are logged to a dedicated audit table with user, timestamp, and details.
Destructive operations	Modifications require explicit user confirmation before execution.
No external dependencies	The tool uses only ServiceNow's built-in APIs and the customer's configured LLM. No external services, no data egress.

Business Value

Quantifiable Benefits

Metric	Current State	With Platform Assistant	Impact
POC setup time	4-8 hours per feature	1-2 hours per feature	60-75% reduction
Troubleshooting time	30-60 min per issue (avg)	5-15 min per issue	50-75% reduction
SC onboarding	2-4 weeks to platform proficiency	Immediate assistance from day 1	Accelerated ramp
AI tool usage on customer POCs	0% (blocked by privacy policy)	100% (privacy-safe)	New capability

Strategic Benefits

1. **Competitive differentiation:** No competitor offers a privacy-safe AI assistant that works within the ServiceNow platform using the customer's own LLM.
 2. **SC productivity at scale:** Every SC benefits immediately — the tool encodes platform expertise that currently lives only in senior SC knowledge.
 3. **Customer confidence:** Demonstrating that AI tools can work within their data governance framework builds trust in our approach to AI-powered services.
 4. **Reusable for customers:** The architecture could be packaged as a deployable solution for customer admins, creating additional value.
-

Delivery Plan

Phase 1: Foundation (2 weeks)

- Deploy the LLM proxy skill on a test instance
- Validate server-side LLM invocation works via GenAI Controller
- Build the core agentic loop with 3 exploration tools (query, schema, record lookup)
- Deliver a working CLI prototype

Milestone: SC can connect to an instance and ask natural language questions about its data.

Phase 2: Diagnostic Tools (2 weeks)

- Add log analysis, ACL checking, script search, flow execution history tools
- Implement the system prompt with ServiceNow platform knowledge
- Add session history and context management

Milestone: SC can diagnose common configuration issues through conversation.

Phase 3: Setup & Configuration Tools (2 weeks)

- Add record creation/update tools with confirmation prompts
- Implement audit logging
- Add one-command setup (sncode-setup)

Milestone: SC can both diagnose and fix issues, with full audit trail.

Phase 4: Polish & Documentation (1 week)

- Error handling, edge cases, timeout management
- User guide and example sessions
- Internal demo and feedback collection

Milestone: Tool ready for broader SC team adoption.

Technical Requirements

Requirement	Detail
ServiceNow version	Yokohama Patch 4+ or Xanadu (for NASK support)
Plugins required	Now Assist, Generative AI Controller, Integration Hub
LLM provider	Any configured via GenAI Controller (Now LLM, BYOK, etc.)
Instance access	Admin role (for setup); itil or equivalent (for usage)
Client requirements	Python 3.6+, terminal access, network access to instance

Instance Footprint

The tool deploys minimal records to the instance:

Records	Count	Purpose
Scripted REST API (definition + endpoints)	4	Chat endpoint, status, tools list
Skills Kit skill (full set)	~14	LLM proxy for GenAI Controller access
Custom table	1	Audit log for modifications

All records are contained and can be removed cleanly.

Success Criteria

1. **Privacy validation:** Confirmed via network trace that zero customer data leaves the instance during tool usage
2. **Time savings:** 50%+ reduction in common troubleshooting tasks (measured in pilot)
3. **Adoption:** 5+ SCs actively using the tool within 30 days of release
4. **Accuracy:** LLM-generated queries and diagnoses are correct 80%+ of the time

Risks & Mitigations

Risk	Likelihood	Impact	Mitigation
GenAI Controller API changes in future ServiceNow releases	Medium	Medium	Abstract the LLM invocation layer; update as needed per release
LLM response quality varies by provider	Medium	Low	Tool works with any model; SCs can switch models if quality is insufficient
Instance performance impact from agentic loop	Low	Medium	Loop limited to 5-8 iterations; uses standard GlideRecordSecure (same as UI usage)
Security review required for Script Runner pattern	Medium	Medium	Use GlideRecordSecure everywhere; audit logging; confirmation for destructive ops

Appendix: Prior Art

This project builds on reverse-engineering work already completed for the ServiceNow AI Agent and Skills Kit frameworks. We have:

- Working automation for creating Skills Kit skills (~24 API calls, fully automated)
- Working automation for creating AI Agent Use Cases (~8 API calls, fully automated)
- Deployed Script Runner API for server-side script execution
- Deployed Provisioning API for Gen AI record creation
- Discovery of GlideUpdateManager2 technique for bypassing ACL-protected operations
- Complete documentation of both frameworks (Skills Kit + AI Agent)

The Platform Assistant is the natural next step — using these capabilities to build an interactive tool rather than one-shot automation scripts.